

GPU Computing

Patterns for massively parallel programming (part 1)

Stencil Pattern and Shared Memory

Edwin Carlinet, Joseph Chazalon

`firstname.lastname@epita.fr`

Fall 2023

EPITA Research Laboratory (LRE)



Non-local Access Pattern with Tiling 🌶️

Tiling in shared memory 🌶️🌶️

Using Shared Memory for Transposition

Bank conflicts 🌶️🌶️🌶️

Conclusion 🌶️


Non-local Access Pattern with Tiling 🌶️

Stencil Pattern

The computation of a single pixel relies on its neighbors

Use case:

- Dilation/Erosion
- Box (Mean) / Convolution Filters
- Bilateral Filter
- Gaussian Filter
- Sobel Filter



x direction

light green	green	dark green
light blue	white	dark blue
medium blue	medium blue	dark blue

 \times

-1	0	1
-2	0	2
-1	0	1

 $=$

y direction

light green	green	dark green
light blue	white	dark blue
medium blue	medium blue	dark blue

 \times

1	2	1
0	0	0
-1	-2	-1

 $=$

Naive Stencil Implementation

Local average with a rectangle of radius r . (Ignoring border problems for now).

```
__global void boxfilter(const int* in, int* out, int w, int h, int r)
{
    int x = blockIdx.x * blockDim.x + threadIdx.x;
    int y = blockIdx.y * blockDim.y + threadIdx.y;
    if (x < r || x >= w - r) return;
    if (y < r || y >= h - r) return;

    int sum = 0;
    for (int kx = -r; kx <= r; ++kx)
        for (int ky = -r; ky <= r; ++ky)
            sum += in[(y+ky) * w + (x+kx)];    // <==== \!/
    out[y * w + x] = sum / ((2*r+1) * (2*r+1));
}
```

Naive Stencil Performance

- Let's say, we have this GPU:
Peak power: 1 500 GFlops and Memory Bandwidth: 200 GB/s
- All threads access global memory
 - 1 Memory access for 1 FP Addition
 - Requires $1\,500 \times \text{sizeof}(\text{float}) = 6 \text{ TB/s}$ of data
 - But only 200 GB/s mem bandwidth \rightarrow 50 GFLOPS (3% of the peak) 📉

- Let's say, we have this GPU:

Peak power: 1 500 GFlops and Memory Bandwidth: 200 GB/s

- All threads access global memory
 - 1 Memory access for 1 FP Addition
 - Requires $1\,500 \times \text{sizeof}(\text{float}) = 6 \text{ TB/s}$ of data
 - But only 200 GB/s mem bandwidth \rightarrow 50 GFLOPS (3% of the peak) 🗣️


⚠️ Many programs are bandwidth-limited and not compute-limited

Compute-to-global-memory-access ratio

We need to have $\#FLOP / \#GlobalMemAccess \geq 30$ to reach the peak

Naive Stencil Performance

163 ms for a 24 MPix image



- Problem: too many access to global memory

- Solution: tiling; copy data to shared memory per block first

Tiling in shared memory 🌶️🌶️

Tiling and memory privatization in shared memory



Figure 1 - Tile subdivision used to render a frame

For each block:

- read the tile from global to private block memory
- process the block
- write the tile from the private block memory to global memory

```
void mykernel() {  
    __shared__ float private_mem[TILE_WIDTH][TILE_WIDTH];  
}
```

Collaborative loading and writing when BLOCKDIM = TILEDIM


- All threads load one or more data
- Access must be **coalesced**
- Use barrier synchronization to make sure that all threads are ready to start the phase

```
void tiledKernel(unsigned char * in, unsigned char * out, int w, int h) {  
    __shared__ float tile[TILE_WIDTH][TILE_WIDTH];  
  
    int x = threadIdx.x + blockDim.x * blockIdx.x;  
    int y = threadIdx.y + blockDim.y * blockIdx.y;  
    // Load  
    if (x < w && y < h)  
        tile[threadIdx.y][threadIdx.x] = in[y * pitch + x];  
    __syncthreads();  
    // Process  
    __syncthreads();  
    // Write  
    if (x < w && y < h)  
        out[y * pitch + x] = tile[threadIdx.y][threadIdx.x];  
}
```

1 thread \leftrightarrow multiple loads


```
{  
    __shared__ float tile[TILE_WIDTH][TILE_WIDTH];  
    int* block_ptr = in + ...; // Compute pointer to the beginning of the tile  
    for (int y = threadIdx.y; y < TILE_WIDTH; y += blockDim.y)  
        for (int x = threadIdx.x; x < TILE_WIDTH; x += blockDim.x)  
        {  
            if (x < width && y < height)  
                tile[y][x] = block_ptr[y * pitch + x];  
        }  
    __syncthreads();  
}
```

Handling Border 🌶️🌶️🌶️



1. Add border to the image to have in-memory access
2. Copy tile + border to shared memory

1. The bad way: each thread copies one value and border threads are then idle.



2. The good way: 1 thread \leftrightarrow multiple loads


```
// TILE_WIDTH = blockDim.x + 2
__shared__ int tile[TILE_WIDTH][TILE_WIDTH]; // Alloc with the size of the block + border

int* block_ptr = in + ...;
for (int i = threadIdx.y; i < TILE_WIDTH; i += blockDim.y)
    for (int j = threadIdx.x; j < TILE_WIDTH; j += blockDim.x)
        data[i][j] = block_ptr[i * pitch + j];

__syncthreads();
```

Stencil Pattern with Tiling Performance

- Global memory: **163 ms** for a 24 MPix image
- Local memory: **116 ms** for a 24 MPix image (30% speed-up)




Using Shared Memory for Transposition

Naive transposition

```
int x = blockIdx.x * blockDim.x + threadIdx.x;
int y = blockIdx.y * blockDim.y + threadIdx.y;

// transpose with boundary test
if (x < w && y < h)
    out[x * width + y] = in[y * width + x]
```

Where are non-coalesced access ?



Naive transposition


```
int x = blockIdx.x * blockDim.x + threadIdx.x;
int y = blockIdx.y * blockDim.y + threadIdx.y;

// transpose with boundary test
if (x < w && y < h)
    out[x * width + y] = in[y * width + x]
```

Where are non-coalesced access ?


→ `a[x][y]`

- Reads are coalesced
- Write are strided



Tiled transposition in shared memory (1/2)

Step 1: Read a row of a block from global memory and write to a row of shared memory.



Shared Memory



Step 2: Read a column from shared memory and write to a row of a block to global memory.



Transposed Matrix


Tiled transposition in shared memory (2/2)

```
void tiledKernel(unsigned char * in, unsigned char * out, int w, int h) {
    __shared__ float tile[TILE_WIDTH][TILE_WIDTH];

    int x = threadIdx.x + blockDim.x * blockIdx.x; // src
    int y = threadIdx.y + blockDim.y * blockIdx.y; // src
    int X = threadIdx.x + blockDim.y * blockIdx.y; // dst
    int Y = threadIdx.y + blockDim.x * blockIdx.x; // dst

    // Load a line
    if (x < w && y < h)
        tile[threadIdx.y][threadIdx.x] = in[y * pitch + x];

    __syncthreads();
    // Write a line from a column in private mem
    if (x < w && y < h)
        out[Y * pitch + X] = tile[threadIdx.x][threadIdx.y];
}
```




Performance (GB/s on TESLA K40)

Copy (baseline)	Transpose Naive	Transpose Tiled
177.15 GB	68.98	116.82

Can we do better ?

Bank conflicts 🌶️🌶️🌶️


About shared memory

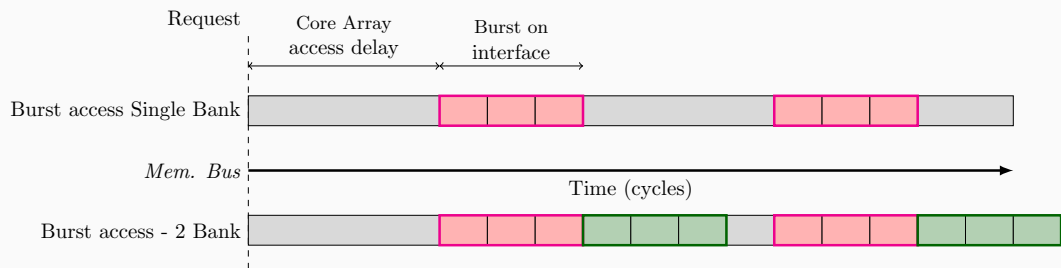


GTX 1080 (Pascal)		Size	Bandwidth	Latency
L1 Cache (per SM)	Low latency	16 or 48K	1,600 GB/s	10-20 cycles
L2 Cache		1-2M		
Global	High latency	8GB	320 GB/s	400-800 cycles

DRAM Banks


- Bursting: access multiple locations of a line in the DRAM core array (horizontal parallelism)
- 2 more forms of parallelism: channels & banks (vertical pipelining)
1 processor has many channels (memory controller) with a bus that connects a set of DRAM banks (core array) to the processor.






Bank conflicts in shared memory

- If 2 threads try to perform 2 **different** loads in the same bank → **Bank conflict**
- Every bank can provide 64 bits every cycle
- Only two modes:
 - Change after 32 bits
 - Change after 64 bits



`load DATA[tid.x]`


No Conflict




`load DATA[42]`

No conflict if loading the same address (broadcast)


- 2-way conflicts



- 2-way conflicts






Conflict = Serialized access (👎)




Concrete Example for Shared Memory

- Bank size: 4B = 4 uint8
- 32 Banks - Many Channels
- Warp Size = 32 threads







op	Items	Bank used	Conflict Free	#Loads
load M[tid.x]				




op	Items	Bank used	Conflict Free	#Loads
load $M[tid.x]$	$[0, 1, \dots, 31]$	$[0, 1, \dots, 7]$	✓	1x8
load $M[tid.x \% 4]$				




op	Items	Bank used	Conflict Free	#Loads
load $M[tid.x]$	$[0, 1, \dots, 31]$	$[0, 1, \dots, 7]$	✓	1x8
load $M[tid.x \% 4]$	$[0, 1, 2, 3, 0, 1, 2, 3, \dots]$	$[0]$	✓	1x1
load $M[tid.x + 1]$				




op	Items	Bank used	Conflict Free	#Loads
load $M[tid.x]$	$[0, 1, \dots, 31]$	$[0, 1, \dots, 7]$	✓	1x8
load $M[tid.x \% 4]$	$[0, 1, 2, 3, 0, 1, 2, 3 \dots]$	$[0]$	✓	1x1
load $M[tid.x + 1]$	$[1, 2, 3, \dots, 32]$	$[0, 1, \dots, 8]$	✓	1x9
load $M[tid.x * 2]$				



op	Items	Bank used	Conflict Free	#Loads
load M[tid.x]	[0,1,...,31]	[0,1,...,7]	✓	1x8
load M[tid.x % 4]	[0,1,2,3,0,1,2,3...]	[0]	✓	1x1
load M[tid.x + 1]	[1,2,3,...32]	[0,1,...,8]	✓	1x9
load M[tid.x * 2]	[0,2,4,...62]	[0,1,...,15]	✓	1x16
load M[tid.x * 8]				




op	Items	Bank used	Conflict Free	#Loads
load M[tid.x]	[0,1,...,31]	[0,1,...,7]	✓	1x8
load M[tid.x % 4]	[0,1,2,3,0,1,2,3...]	[0]	✓	1x1
load M[tid.x + 1]	[1,2,3,...32]	[0,1,...,8]	✓	1x9
load M[tid.x * 2]	[0,2,4,...62]	[0,1,...,15]	✓	1x16
load M[tid.x * 8]	[0,8,...,248]	[0,2,...,30]	✗	2x16
load M[tid.x * 12]				



op	Items	Bank used	Conflict Free	#Loads
load $M[tid.x]$	$[0, 1, \dots, 31]$	$[0, 1, \dots, 7]$	✓	1x8
load $M[tid.x \% 4]$	$[0, 1, 2, 3, 0, 1, 2, 3 \dots]$	$[0]$	✓	1x1
load $M[tid.x + 1]$	$[1, 2, 3, \dots, 32]$	$[0, 1, \dots, 8]$	✓	1x9
load $M[tid.x * 2]$	$[0, 2, 4, \dots, 62]$	$[0, 1, \dots, 15]$	✓	1x16
load $M[tid.x * 8]$	$[0, 8, \dots, 248]$	$[0, 2, \dots, 30]$	✗	2x16
load $M[tid.x * 12]$	$[0, 12, \dots, 372]$	$[0, 1, \dots, 31]$	✓	1x32

Bank conflicts in Transpose



```
__shared__ a[16][16];
```

```
Y = by*16+ty;
```


```
X = bx*16+tx;
```

```
a[y][x] = A[Y][X]
```

```
_syncthreads()
```

```
B[Y][X] = a[x][y];
```

Bank conflicts in Transpose



```
__shared__ a[16][16];
```

```
Y = by*16+ty;
```

```
X = bx*16+tx;
```

```
a[y][x] = A[Y][X]
```

```
_syncthreads()
```

```
B[Y][X] = a[x][y];
```

Reading a column may create bank conflicts

Solution to bank conflicts

- With padding (to $\text{WRAP_SIZE} + 1$)

```
__shared__ a[17][17];
```


```
Y = by*16+ty;
```

```
X = bx*16+tx;
```

```
a[y][x] = A[Y][X]
```

```
_syncthreads()
```

```
B[Y][X] = a[x][y];
```



row & column access pattern

- Index mapping function

$$f: (x,y) \rightarrow y * 16 + (x+y) \% 16$$

```
__shared__ a[...];
```



```
Y = by*16+ty;
```

```
X = bx*16+tx;
```

```
a[f(x,y)] = A[Y][X]
```

```
__syncthreads()
```

```
B[Y][X] = a[f(y,x)]
```



row & column access pattern

Performance (GB/s on TESLA K40)

Copy (baseline)	Transpose Naive	Transpose Tiled	Transpose Tiled+Pad
177.15 GB	68.98	116.82	121.83

Conclusion 🌶️

Shared memory (Summary)

- Superfast access (almost as fast as registers) 👍
- Useful for block-wise collaborative computation (next course) 👍
- But limited resources (64~96Kb by SM) 👎

Use it carefully to avoid reducing the occupancy...

Number of active warps divided by the maximum number of warps that could be executed on the SM.

Generation	Warps per SM	Warps per scheduler	Active threads limits
Maxwell (5.2)	64	16	2048
Pascal (6.1)	64	16	2048
Volta (7.0)	64	16	1024
Turing (7.5)	32	8	1024

If Shared Memory usage ↗, then the number of ACTIVE Warp / SM ↘