

GPU Computing

Edwin Carlinet, Joseph Chazalon

`firstname.lastname@epita.fr`

Fall 2024

EPITA Research Laboratory (LRE)




GPU vs CPU architectures 🌶️

How to explain that:


- CPU are low-latency
- GPU are high-throughput

It's all about data... the CPU:




- Optimized for low-latency **access** (many memory caches)
- Control logic for out-of-order and speculative execution

Intel i7




IBM Power 8 (2014)




It's all about data... the GPU:

CPU



- Low-latency **access**
- Many control logic

GPU




- Throughput computation (ALUs)
- Tolerant to memory latency

But how... ?

Hiding latency With thread parallelism & pipelining

So... you want to hide the latency of getting data from global memory... how ?


1 CPU Core




Hiding latency With thread parallelism & pipelining

So... you want to hide the latency of getting data from global memory... how ?

1 CPU Core



1 GPU SMP (Streaming Multiprocessor)



CPU

GPU

- low-latency memory to get data ready

- memory latency hidden by **pipelining**

Latency hiding


- = do other operations when waiting for data
- = having a lot of parallelism
- = having a lot of data
- will run faster
- but not faster than the peak
- what is the peak by the way ?

It's all about data... Little's law

- Customer arrival rate: **Throughput**
- Customer time spent: **Latency**
- Average customer count: **Concurrency (Data in the pipe)" = throughput * Latency**

Concurrency is the number of items processed at the same time.

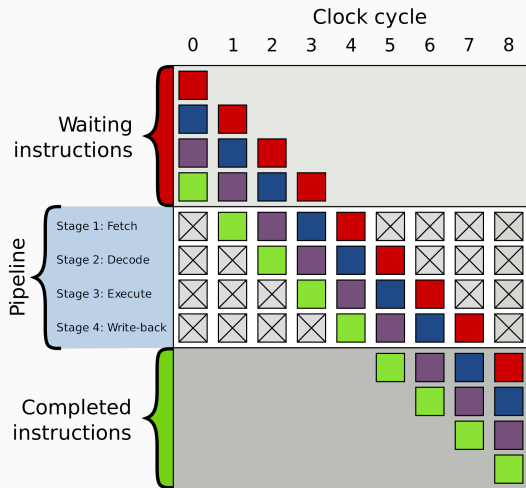
	Latency	Peak Throughput	Needed Concurrency
GPU-arithmetic	24 cycles	8 IPC	192 inst
GPU-memory	350 ns	190 GB/s	65K



Hiding latency

With thread parallelism & pipelining

Note that pipelining exists on CPUs (cycle de Von Neumann):




- IF instruction fetch
- ID instruction decode
- OPF Operand fetch
- EX execute
- WB result write back

Pipeline at **Instruction Level** vs pipeline at **Thread (Warp) Level**

More about forms of parallelism (the why!) 🌶️🌶️


Vertical parallelism for latency hiding

Pipelining keeps units busy when waiting for dependencies, memory



Horizontal parallelism for throughput

More units working in parallel



More about forms of parallelism (the how!) 🌶️🌶️

Instruction-Level Parallelism (ILP)

Between independent instructions.

1. add **r3** ← r1, r2
2. mul **r0** ← r0, r1
3. sub r1 ← **r3**, **r0**

'1 and '2 run concurrently

More about forms of parallelism (the how!) 🌶️🌶️

Instruction-Level Parallelism (ILP)

Between independent instructions.

1. add **r3** ← r1, r2
2. mul **r0** ← r0, r1
3. sub r1 ← **r3**, **r0**

Thread-Level Parallelism (TLP)

Between independent execution contexts: threads


'1 and '2 run concurrently

Thread 1



add

Thread 2



mul

More about forms of parallelism (the how!) 🌶️🌶️

Instruction-Level Parallelism (ILP)

Between independent instructions.

1. add **r3** \leftarrow r1, r2
2. mul **r0** \leftarrow r0, r1
3. sub r1 \leftarrow **r3**, **r0**

Thread-Level Parallelism (TLP)

Between independent execution contexts: threads


'1 and '2 run concurrently

Thread 1



add

Thread 2



mul

Data-Level Parallelism (DLP)

Between elements of a vector: same operation on several elements

vadd r \leftarrow a, b

a₁ a₂ a₃

+

b₁ b₂ b₃

r₁ r₂ r₃

Extracting parallelism 🌶️🌶️

	Horizontal	Vertical
ILP	Superscalar	Pipeline
TLP	Multi-cores / SMT	Interleaved / Switch-on-event multi-threading
DLP	SIMD / SIMT	

Parallel architectures & parallelism 🌶️🌶️

CPU (Intel Haswell)

	Hor.	Vert.
ILP	8	✓
TLP	4	2
DLP	8	

- 8 ALUs for executing non-dependent instructions
- 4 cores. Physical cores
- 4*2 logical hyper-cores
- Lane of 8x32bits SIMD registers supporting AVX 256

General-purpose multi-cores: balance ILP, TLP and DLP

Parallel architectures & parallelism 🌶️🌶️

CPU (Intel Haswell)

	Hor.	Vert.
ILP	8	✓
TLP	4	2
DLP	8	

- 8 ALUs for executing non-dependent instructions
- 4 cores. Physical cores
- 4*2 logical hyper-cores
- Lane of 8x32bits SIMD registers supporting AVX 256

General-purpose multi-cores: balance ILP, TLP and DLP

GPU (Nvidia Kepler)

	Hor.	Vert.
ILP	2	
TLP	16x4	64
DLP	32	

- Dual instruction issue for executing non-dependent instructions
- 16 Multiprocessors (physical cores) Can execute 4 simultaneous warps
- Multithreading (64 warps / SM)
- 128 (4 x 32) CUDA cores. SIMT of width 32

GPU: focus on DLP, TLP horizontal and vertical.

- All processors use hardware to turn parallelism into performance
- GPUs focus on Thread-level and Data-level parallelism